

Arizona Flash Flood Inundation AI Early Warning System

(AFFI)

A Physics-Guided Artificial Intelligence Framework for
Real-Time Flood Inundation Forecasting in Rural Southern Arizona

May 11, 2026 — original. **Updated with dated August 2026 annotations** (this document; see orange UPDATE boxes throughout — original text is unchanged).

Prepared by: Solman Raju Sarva, MSc — AI & Hydrology Researcher and William D. O'Brien, PE, CFM — Principal Engineer, NextGen Engineering Inc.

Pilot Watershed: Upper Sonoita Creek Watershed, near Patagonia, Santa Cruz County, AZ

Training Watershed (as printed May 11): Walnut Gulch, near Tombstone, Cochise County AZ

UPDATE — VERIFIED 2026-08-10

The training watershed printed above is not what the deployed model actually trains on. Verified directly against `config/task2.yaml` and file timestamps in `models/`:

- The anchor/training basin was switched from Walnut Gulch to the **Babocomari River** (real USGS gauge 09471000, Cochise/Santa Cruz County AZ) on **2026-05-27** — `config/task2.yaml`'s own header states why: *"swapped Walnut Gulch (fake ARS) → Babocomari River (real USGS 09471000)"*. The original Walnut Gulch dataset used at the time of this May 11 draft turned out to be synthetic/placeholder ("fake ARS") data, not the real USDA ARS record.
- The deployed model artifacts confirm this: `models/xgb_magnitude.joblib` is dated **2026-05-27**; `models/classifier_best.pt` and `models/best_inference_config.json` (the final event-gate model and its config) are dated **2026-06-09**.
- **Every mention of "Walnut Gulch" as the anchor watershed later in this document (Sections 1, 4.4, 6.1, Table 5) is therefore out of date as of 2026-05-27** — left in place below exactly as originally written, with an UPDATE box at each occurrence.

SCOPE: This "white paper" presents a technical framework for Flood Inundation Early Warning pilot tested on the **Upper Sonoita Creek Watershed (HUC 150503010204)** in Santa Cruz County, Arizona. The architecture, methodology, and deliverables described herein are designed to be **transferable to any watershed in Arizona or the United States** with minimal reconfiguration. This "white paper" is aimed at building collaboration for the full development of AFFI early warning systems.

1. Executive Summary

Flash floods are the deadliest natural hazard in Arizona, killing an average of 14 people per year and causing \$180 million in annual property damage. The fundamental problem is not the flood — it is the **time delay from storm forecasts to seeing where to warn people**. Current physics-based flood simulation tools require 4 to 6 hours to produce a flood inundation map. A monsoon flash flood in Southern Arizona can arrive in under 2 hours. That gap between warning and flood impact is where lives are lost.

The Arizona Flash Flood Inundation AI Early Warning System (AFFI) closes this gap. AFFI is a six-staged physics-guided artificial intelligence pipeline that ingests real-time weather forecast data and delivers a probabilistic flood inundation map to authorized government personnel at their phone or Emergency Operations Center (EOC) in about 60 seconds — 360 times faster than current methods — with accuracy comparable to physics-based models. Table ES-1 outlines the benefits of AFFI.

This document describes the complete technical architecture of AFFI, piloted on the **Upper Sonoita Creek Watershed (HUC 150503010204)** in Patagonia, Santa Cruz County, Arizona. Every component of this system is designed for **transferability**: the same AFFI System, with watershed-specific data substitution, can be deployed for any watershed in Arizona or the United States.

To improve model robustness and transferability, the AI system is first trained on USDA Walnut Gulch Experimental Watershed hydrologic and rainfall-runoff data and then fine-tuned for the targeted watershed deployment area. (Since the late 1950s the Walnut Gulch Watershed has long-term hydrology, erosion, and water quality data collection and research programs near Tombstone, Az.) The transfer-learning approach allows AFFI to leverage high-quality historical flood dynamics from a well-instrumented watershed while adapting to the terrain, hydrology, and storm-response characteristics of the specific target watershed.

UPDATE — 2026-08-10 **Anchor watershed (repeats the cover-page correction):** the paragraph above still names Walnut Gulch. As of 2026-05-27 the real anchor is the **Babocomari River** (USGS 09471000, real gauge record), not Walnut Gulch — see the cover-page UPDATE box for the full explanation. Everything else in this paragraph (the transfer-learning strategy itself: train on a data-rich anchor, then fine-tune on the pilot watershed) is still accurate and still how the system works today.

The full development of the AFFI system is broken into **tasks**, each with specific, verifiable deliverables. All data sources are publicly available from federal agencies at no licensing cost. The flood plain modeling software, Hec-Ras, is publicly available. Task 1 has already been completed and validated.

UPDATE — 2026-08-10 **Task status, three months on:** as of this update, Tasks 1, 2, 4, and 6 are implemented and running live (not just Task 1). Task 3 is delivered in a defensible interim form — an analytical FEMA/USGS flood-depth library, not the U-Net described later in this document — with the neural surrogate implemented as a parallel, not-yet-operational track. Full detail: docs/AFFI_whitepaper_2026-08-10.md, Section 5.

Table ES-1: Benefits of AFFI

Metric	Current Capability	AFFI	Improvement
Time to flood map	4–6 hours (HEC-RAS)	< 60 seconds	360× faster

Watersheds covered	Gauged only (~12% of AZ)	Any watershed in AZ	Statewide Applicability
Forecast lead time	0–2 hours	Up to 7 days	84× longer
Ensemble members	1 (deterministic)	31 (probabilistic)	Uncertainty quantified
Cost per forecast	~\$2,400 (EOC activation)	~\$0.04 (cloud compute)	60,000× cheaper
Ungauged basin coverage	Requires stream-gauge calibration	Any Arizona watershed	Ungauged-ready

UPDATE — 2026-08-10 **"31 (probabilistic)" ensemble members:** the GFS weather forecast genuinely has 31 members (this is still accurate, unchanged). What changed: the pipeline does not run all 31 independently through hydrology + hydraulics. It propagates the **P10/P50/P90 rainfall percentiles** (already derived from the 31-member ensemble) through the pipeline to produce a best/likely/worst triplet — faster, and honest about what the live forecast packet actually carries. Full 31-member independent propagation is a stated upgrade path, not yet built. See Section 4.6's update box.

2. Problem Statement: Flash Flooding in Rural Southern Arizona

2.1 Human Cost

Between 2000 and 2023, flash floods killed 287 people in Arizona — more than tornadoes, hurricanes, and earthquakes combined in the same period (NOAA Storm Data, 2023). Most fatalities occurred in vehicles, where drivers encountered flooded roadways with little or no advance warning.

The pilot area for AFFI is the Upper Sonoita Creek watershed near Patagonia in Santa Cruz County, Arizona (HUC 150503010204). This drains 55.4 mi² (143.6 km²) of semi-arid terrain and has experienced 12 major flood events since 1980, with peak flows exceeding the 50-year return period on three occasions. State Route 82 near Patagonia — the only evacuation route for 3,200 people — has been rendered impassable by flooding on multiple occasions with little or no advance warning to emergency managers.

Example: The 2006 Santa Cruz County Monsoon Flood: A single monsoon event caused \$12 million in infrastructure damage, displaced 340 residents, and closed State Route 82 for 72 hours. Emergency managers received no advance flood inundation map. Response was entirely reactive. This event is the baseline scenario AFFI is designed to mitigate.

UPDATE — 2026-08-10 The 55.4 mi² figure: confirmed correct for the HUC-12 boundary and unchanged — but it's a *different, smaller* spatial extent than the ~197–209 mi² (~510 km²) figure used elsewhere in this system (config, hydrology scaling) for the full USGS gauge 09481500 contributing drainage area. Both are real; verified this session against the gauge's own published drainage area (~209 mi²). The gauge sees a larger upstream area than just this HUC-12.

2.2 Infrastructure and Economic Cost

The economic cost of flash flooding in Arizona averages \$180 million annually in direct damage to roads, bridges, utilities, and structures (FEMA, 2022). Indirect costs include emergency-response mobilization (~\$2,400/hour for a county EOC), post-flood debris removal (\$180,000–\$450,000 per major event), long-term infrastructure repair (\$1.2M–\$8M per bridge or road segment), and agricultural losses (\$15M–\$40M per severe monsoon season in Southern Arizona).

2.3 The Technology Gap

Southern Arizona's rural flash flood problem is uniquely difficult. The region is characterized by hundreds of ephemeral washes and arroyos — channels that are dry for 300+ days per year but can carry catastrophic flows within minutes of drenching by a monsoon rain cell. These channels are ungauged, ephemeral, braided, and peak with a rapid response to rainfall. The time of concentration as short as 15–30 minutes for small sub-watersheds.

Current flood mapping operational tools — NOAA's National Water Model, FEMA Flood Insurance Rate Maps, Google Flood Hub — were designed for large, gauged river systems. They provide valuable national-scale flood information but do not address the specific challenge of rural, ungauged, ephemeral flash flood prediction in Southern Arizona. **No existing operational system provides real-time, event-specific flood inundation maps for ungauged rural arroyos in Southern Arizona with lead times sufficient for emergency management action. AFFI fills this gap.**

2.4 The Climate Trend

Climate projections for the Southwest United States indicate a 15–40% increase in extreme precipitation intensity during monsoon events by 2050 (IPCC AR6, 2021). Flood events that currently occur once every 50 years will occur once every 20–30 years by mid-century. The window for investing in flood inundation early warning infrastructure — before these events become routine — is urgent.

3. Existing Systems and Their Limitations

A thorough review of existing flood inundation forecasting systems confirms that no current tool adequately addresses the rural Southern Arizona flash flood problem. The following analysis is not a criticism of these systems — each is excellent for its intended purpose. It is a demonstration that a gap exists which none of them fills.

Table 1: Comparison of AFFI with other Flood Inundation Mapping Systems

System	Forecast Type	Lead Time	Ungauged Basins	Rural Arroyos	Real-Time Maps
AFFI (Proposed)	Probabilistic inundation maps	Up to 7 days	Yes	Yes	Yes
NOAA NWM	Streamflow forecast	Up to 7 days	Limited	No	No
Google Flood Hub	Flood forecast + maps	Up to 7 days	No	No	Partial
FEMA FIRM	Static hazard zones	None (static)	Partial	Partial	No
First Street Foundation	Risk scoring	None (static)	Partial	No	No
FLASH (RoyalHaskoningDHV)	Short-lead flood alerts	Hours to 1 day	No	No	Partial

UPDATE — 2026-08-10 "AFFI (Proposed)": AFFI is no longer proposed-only — it's a running, deployed pilot as of this update, with a live dashboard, scheduler, and government API. The row label above is left as originally written; treat it as "AFFI (Pilot — validated)."

3.1 Detailed Limitations of Each System

NOAA National Water Model (NWM): The outstanding national-scale streamflow forecasting system. However, it depends on stream-gauge calibration data, is not designed for ephemeral channels, and produces streamflow forecasts rather than spatial flood-inundation maps. Reference: NOAA Office of Water Prediction.

Google Flood Hub: Provides flood forecasting for gauged rivers globally. But does not cover ungauged basins or ephemeral arroyos. Coverage in Arizona is limited to major perennial rivers. Reference: Google Flood Hub.

FEMA Flood Insurance Rate Maps (FIRM): Static maps showing 100-year and 500-year flood zones. Not event-specific, not real-time, and not updated frequently enough to reflect current channel or specific storm conditions. Reference: FEMA Flood Map Service Center.

First Street Foundation: Provides property-level flood risk scores. Static, not event-specific, and not designed for emergency operations. Reference: First Street Foundation.

FLASH System: Provides short-lead flash flood alerts based on radar rainfall. Does not produce spatial inundation maps and does not cover ungauged basins with sufficient resolution for rural Arizona. Reference: NOAA Flash Flood Guidance.

4. Technical Solution: The AFFI Architecture

4.1 System Overview — The Physical Pipeline

AFFI follows the physically correct sequence for flood prediction. This sequence mirrors nature, how rainfall becomes a flood:


Rainfall Forecast → Mean Areal Watershed Precipitation → Runoff / Discharge → Flood Depth & Extent → Probability Maps → Alert

Each arrow represents a distinct physical transformation. Each transformation is handled by the most appropriate computational method — physics-based where physics is well understood, AI-based where AI has been proven to match or exceed physics-based accuracy at a fraction of the computational cost.

Table 2: Six Stages of AFFI Architecture

Stage	Input	Output	Method	Time
Stage 1: Meteorology	GFS/HRRR ensemble forecast	Mean Areal Precipitation (MAP) per watershed	Statistical downscaling	~5 sec
Stage 2: Hydrology	MAP time series + watershed attributes	Peak discharge (ft ³ /s) + hydrograph	Physics-guided LSTM	~8 sec
Stage 3: Hydraulics	Peak discharge + terrain (DEM)	Flood depth + extent raster	U-Net / ResUNet	~45 sec
Stage 4: Probabilistic Output	31 ensemble flood maps	Probability maps, uncertainty maps	Ensemble statistics	~2 sec
Stage 5: Benchmarking	Predicted discharge	Return period classification (2yr–100yr)	NOAA Atlas 14 comparison	~1 sec
Stage 6: Alert	All above outputs	Alert packet (JSON) + dashboard	Rule-based logic + API	~1 sec

Total pipeline run time: ~60 seconds. This compares to 4–6 hours for the equivalent hydrology (HEC-HMS model or equal) + hydraulic (HEC-RAS model or equal) physics-based simulation chain.



4.2 Why This Architecture and Not Others

After reviewing over 200 published AI flood prediction systems, this six-stage architecture is the only design that satisfies all five operational requirements simultaneously:

- **Speed:** ~60 seconds end-to-end — required for flash flood warning
- **Accuracy:** Comparable to physics-based models — required for government trust
- **Transferability:** Works on ungauged watersheds — required for statewide and rural deployment
- **Interpretability:** Each stage has physical meaning — required for regulatory acceptance
- **Uncertainty quantification:** Probabilistic output — required for risk-based decisions

A single end-to-end deep learning model (rainfall directly to flood map) fails on transferability — it only memorizes the training watershed and cannot transfer. A pure physics model fails on speed. The six-stage hybrid presented is the only architecture that passes all five tests.

4.3 Stage 1: Meteorological Interface — Explained

What it does: Ingests ensemble rainfall forecasts from NOAA's Global Forecast System (GFS, 31 members, 6-hourly updates) and the High-Resolution Rapid Refresh (HREF - High-Resolution Ensemble Forecast, ~10 members, hourly, CONUS). For each watershed, computes Mean Areal Precipitation (MAP) is the spatially averaged rainfall over the watershed extent — for forecast horizons of 1, 3, 6, 12, 24, 48, 72, and 168 hours.

Why it matters: A point rainfall measurement tells you how much rain fell at one location. A watershed responds to the average rainfall over its entire area. MAP is the physically correct input to a hydrologic model. Using point rainfall instead of MAP is a common error that leads to significant over- or under-prediction of flood response.

What is AI at this stage? Stage 1 is NOT AI. It is physics-based numerical weather prediction (NWP) from NOAA's operational models. The AI begins in Stage 2. This distinction is important for scientific integrity and regulatory acceptance.

Comparison to 10-year storm rainfall: Each MAP value is compared against NOAA Atlas 14 precipitation frequency benchmarks for Santa Cruz County, producing a return-period classification (2-year through 100-year) and an initial alert level (Advisory / Watch / Warning). This allows emergency managers to immediately understand whether an incoming storm is a routine event or a rare, high-consequence event.

UPDATE — 2026-08-10 **Stage 1 status:** unchanged in method, and this is the one stage where the original description still matches the live system almost exactly. Live, running on a 6-hour automated schedule, producing a real alert packet consumed by the dashboard. Gauge identity corrected — see Table 5's update.

4.4 Stage 2: Physics-Guided LSTM Hydrology Model — Explained

What it does: Converts the MAP time series into a predicted streamflow hydrograph (discharge over time, in ft³/s) using a Long Short-Term Memory (LSTM) neural network trained on Walnut Gulch 70-year anchor dataset, fine-tuned on pilot watershed.

Why LSTM? Rainfall-runoff is a time-dependent process — today's soil moisture depends on yesterday's rainfall. LSTM networks are specifically designed to learn these temporal dependencies. They have been independently

validated on 531 US watersheds by Kratzert et al. (2019), achieving Nash-Sutcliffe Efficiency (NSE) > 0.80 — matching or exceeding calibrated physics-based models.

HOW: LSTM is trained in two stages: base training on a high-quality anchor watershed (Walnut Gulch), followed by fine-tuning on the pilot watershed. (Reference transfer learning in hydrology, Kratzert et al. 2019 already supports this.)

Why "physics-guided"? The LSTM incorporates static watershed attributes (drainage area, slope, soil permeability, vegetation index, elevation, aridity index) from the USGS GAGES-II dataset. These attributes encode the physical characteristics of the watershed into the model, enabling it to predict discharge for watersheds it has never seen — the key to generalizability. Model training follows a two-stage transfer learning strategy. In Stage 1, the LSTM is trained on the Walnut Gulch Experimental Watershed — the world's most complete semi-arid flash flood dataset with 70+ years of record — learning fundamental desert hydrology: rapid infiltration, Hortonian overland flow, ephemeral channel routing, and monsoon precipitation-runoff dynamics. In Stage 2, the base model is fine-tuned on local gauge data for the target watershed (for the Sonoita Creek pilot using only the final layers), requiring as few as 10–15 years of local data. This approach — base training on a data-rich anchor watershed followed by targeted fine-tuning — is established in the hydrology AI literature (Kratzert et al. 2019) and is the mechanism that enables generalization to ungauged watersheds.

Comparison to 10-year storm peak runoff: The predicted peak discharge is compared against the 10-year, 25-year, 50-year, and 100-year return period flows derived from USGS regional regression equations for Arizona. This tells emergency managers "how rare or how extreme" this event is.

UPDATE — 2026-08-10 — the biggest change in this document This entire section describes a single LSTM regressor

trained on Walnut Gulch. Neither claim is what's actually running.

1. **Anchor watershed:** Babocomari River (USGS 09471000), not Walnut Gulch — switched 2026-05-27 because the Walnut Gulch data available at the time was synthetic ("fake ARS"), not real. See the cover-page UPDATE box for the full citation trail.
2. **Architecture:** not a single LSTM regressing a hydrograph. The deployed model is a two-part **hurdle model** — see Figure B below.




Figure B — the actual Stage 2 architecture. The XGBoost regressor only runs on days the LSTM gate fires ($P \geq 0.85$); most days it doesn't fire and discharge is reported as zero. This is why it's called a "hurdle" model.

Why a hurdle model, not one network? Desert channels are dry ~95% of the time and produce almost all the interesting hydrology on a handful of days a year. Splitting *whether* it floods from *how big* matches the physical structure of the problem better than one network trying to do both.

Real held-out test metrics, Upper Sonoita Creek (verified against saved model output 2026-08-10)

Metric	Value
AUC-ROC (event detection)	0.959
AUC-PR (rare-event detection)	0.643
F1 (at $P \geq 0.85$)	0.611
NSE (magnitude)	0.348 — below the 0.75 target; expected for this regime, see Section 7
PBIAS	-2.9%

4.5 Stage 3: U-Net Hydraulic Surrogate Model — Explained

What it does: Converts predicted discharge into a spatial flood inundation map (flood depth and extent raster) using a U-Net convolutional neural network trained on 250 physics-based HEC-RAS flood simulations.

Why U-Net? U-Net was originally developed for medical image segmentation and is now the standard architecture for spatial prediction problems. It takes a 2D input (terrain + discharge) and produces a 2D output (flood depth map) — exactly the transformation needed for hydraulic modeling.

Why train on HEC-RAS? HEC-RAS is the industry-standard open-source hydraulic model used by FEMA and USACE. By training the U-Net on HEC-RAS outputs, AFFI inherits the physical accuracy of HEC-RAS while operating 360× faster. The U-Net is a surrogate (emulator) of HEC-RAS, not a replacement.

Physics constraints enforced: Three physics constraints are applied to the U-Net output: (1) Gravity constraint — water cannot flow uphill; (2) Monotonicity constraint — higher discharge must produce equal or greater flood inundation area and extent; (3) Mass conservation — total predicted runoff volume cannot exceed total rainfall volume. These ensure physical plausibility even for scenarios outside the training distribution.

Comparison to 10-year to 100-year flood inundation: The predicted flood map is compared against pre-computed HEC-RAS flood maps for the 10-year, 25-year, 50-year, and 100-year period events. This provides emergency managers with a direct, intuitive benchmark: "This event is expected to produce flooding comparable to a 25-year storm."

UPDATE — 2026-08-10 **Not what's driving today's maps.** The U-Net/ResUNet described above is real and implemented in code, but it is **not** the source of today's live dashboard maps — it's a parallel development track for when a fuller HEC-RAS simulation bank exists. What actually generates the operational depth maps today is an **analytical FEMA/USGS flood library**: real FEMA regulatory floodplain geometry + FEMA Base Flood Elevations (429 sample points) minus the real USGS 3DEP 10-m DEM, scaled to other return periods via a Leopold hydraulic-geometry relationship anchored to USGS statistical discharge estimates. It is a deliberate, disclosed simplification for the pilot stage — full, defensible, and traceable to named federal sources, but not "250 HEC-RAS simulations" as described above. When a HEC-RAS bank is built out, the same U-Net data structure/interface is reused unchanged.

4.6 Stage 4: Probabilistic Output

The full AFFI pipeline is executed independently for each of the 31 GFS ensemble members, producing 31 flood inundation maps for each forecast event. These maps are stacked and 4 critical flood maps are produced. This ensemble approach quantifies forecast uncertainty and produces the following output products as described in Table 3:

Table 3: Useable Probabilistic Output

Output Product	Description	Primary User
Median flood map	Best-estimate flood depth and extent	EOC Staff
90th percentile map	Worst-case scenario (10% chance of exceeding)	EOC staff, Public Works Director
Probability map	P(depth > 0.5m) for each pixel, 0–100%	EOC Staff
Uncertainty map	Standard deviation of depth across 31 maps	Technical reviewer
Alert classification	Advisory / Watch / Warning vs. return period	County OEM
Time-to-peak estimate	Hours until maximum flooding at each location	Road closure coordinator

UPDATE — 2026-08-10 **"Executed independently for each of the 31 GFS ensemble members" — not what runs today.**

The pilot instead propagates the **P10/P50/P90 rainfall percentiles** (already derived from the 31-member ensemble) through an SCS Curve-Number runoff model + flood-library lookup, producing a best/likely/worst triplet rather than 31 independent maps. This is faster and honest about what the live forecast packet carries (percentiles, not per-member lag-feature histories). The output products in Table 3 above are still real and delivered — median/likely map, worst-case map, life-safety probability map, uncertainty map, alert classification, time-to-peak — just computed from 3 scenarios, not 31. Full 31-member propagation is the stated upgrade path.

4.7 Stage 5: Benchmarking — Explained

What it does: Compares every AFFI-predicted rainfall, discharge, and flood map against established federal engineering benchmarks — NOAA Atlas 14 precipitation frequency data and USGS regional regression equations —

to classify each event by return period (2-year through 100-year). This gives emergency managers an immediate, intuitive answer: "Is this a routine storm or a once-in-50-years event?"

Why it matters: Emergency managers are trained to act on return period language. A raw discharge number (e.g., 4,200 ft³/s) is meaningless without context. Saying "this event is comparable to a 25-year storm" is immediately actionable. Benchmarking translates AI output into the engineering language that government decision-makers trust and are legally required to use.

What is AI at this stage? Stage 5 is NOT AI. It is a deterministic comparison engine using published federal data. This reinforces scientific integrity and regulatory credibility.

4.8 Stage 6: Alert Policy - Explained

The alerts from the AFFI need to support the many other data and real-time voices that the Emergency personnel and public officials must deal with. The following Alert Policy is recommended:

Critical Policy Decision: All AFFI alerts are delivered exclusively to authorized government personnel — county emergency managers, public works directors, and EOC staff. No alerts are broadcast directly to the public. This design eliminates the liability risk of automated public warnings and ensures that all public communications are reviewed and authorized by responsible government officials.

4.9 Lead-Time-Dependent Confidence Framework

As the time of the event approaches, the confidence of the prediction improves. Table 4 shows the rough measure of confidence as the storm approaches.


Table 4: Confidence of Predicted Flooding and Related Actions

Lead Time	Primary Data Source	Confidence Level	Primary Use
2–7 days	GFS ensemble (31 members)	Low–Medium	Pre-positioning resources, early briefings
6–24 hours	HRRR + updated GFS	Medium–High	Road closure decisions, evacuation planning
0–6 hours	MRMS radar + nowcast	High	Active road closures, emergency response

UPDATE — 2026-08-10 Stage 6, delivered — far beyond what this section describes. The alert policy above is exactly what's built and running. What's not captured here: the delivered dashboard is a full React + TypeScript / FastAPI application (Section 6's original plan called for Streamlit/Dash) with a Decision Cockpit (time-to-peak, life-safety probability, the Table-4 confidence framework live), an interactive map with Google-Flood-Hub-style gauge status badges and a flood-probability heatmap, an NWS-style bulletin generator citing Arizona's Stupid Motorist Law (ARS 28-910), a data-staleness watchdog, and a 6-hourly automated scheduler. Full detail in Section 6's update box below and in docs / D6.1_alert_logic_specification.md.

5. Development of AFFI – Detailed Tasks Descriptions with Deliverables

Tasks T1 through T6



Task 1: Meteorological Forecast Interface

Objective: Build and validate the system that ingests real-time weather forecast data and computes Mean Areal Precipitation (MAP) for the Upper Sonoita Creek Watershed, with alert threshold logic benchmarked against NOAA Atlas 14 return periods.

Step-by-Step:

- **Step 1.1 — Define Area of Interest (AOI):** Load the Upper Sonoita Creek Watershed boundary (HUC 150503010204) from the USGS National Map. Define bounding box (N=31.85, S=31.47, E=-110.50, W=-110.90). This is the spatial domain for all precipitation calculations.
- **Step 1.2 — Ingest GFS Ensemble Forecast:** Connect to the Open-Meteo API (free, no key required) to retrieve 31-member GFS ensemble precipitation forecasts at 6-hourly intervals for 7-day lead times. Parse and store as structured arrays.
- **Step 1.3 — Compute Mean Areal Precipitation (MAP):** Spatially average the gridded GFS precipitation over the watershed boundary using area-weighted averaging. Compute MAP for accumulation windows of 1, 3, 6, 12, 24, 48, 72, and 168 hours.
- **Step 1.4 — Load NOAA Atlas 14 Benchmarks:** Load precipitation frequency estimates for Santa Cruz County, AZ from NOAA Atlas 14 Volume 1 for return periods of 2, 5, 10, 25, 50, and 100 years at all accumulation durations.
- **Step 1.5 — Compute Return Period Classification:** For each MAP value and accumulation duration, determine which return period bracket it falls into (e.g., "between 10-year and 25-year"). Compute a Storm Severity Index ($SSI = MAP / MAP_{10yr}$) for each duration.
- **Step 1.6 — Apply Alert Threshold Logic:** Apply the following thresholds calibrated for the Sonoita Creek watershed: Advisory (0.25"/1hr, 0.75"/24hr), Watch (0.50"/1hr, 1.25"/24hr), Warning (1.00"/1hr, 2.00"/24hr). Generate an alert level (GREEN / ADVISORY / WATCH / WARNING) for each ensemble member and lead time.
- **Step 1.7 — Generate Alert Packet:** Compile all outputs into a structured JSON alert packet containing: alert level, MAP values, return period classifications, ensemble statistics (mean, 10th, 50th, 90th percentile), and time-to-peak estimate.
- **Step 1.8 — Dummy Rainfall Stress Test:** Validate the alert logic using synthetic rainfall scenarios spanning the full range from light shower to extreme cloudburst. Confirm that each scenario triggers the correct alert level and return period classification.

Deliverables: D1.1 — Validated Python module: task1_met_interface.py — ingests GFS ensemble, computes MAP, classifies return period, generates alert level.

D1.2 — Alert Packet JSON: task1_alert_packet.json — structured output for all 31 ensemble members, all lead times, all accumulation durations.

D1.3 — Forecast Dashboard: task1_forecast_dashboard.png — 4-panel visualization: MAP time series, ensemble spread, return period comparison, alert level timeline.

D1.4 — Dummy Stress Test Report: task1_dummy_benchmark.png — 2-panel chart showing all synthetic scenarios plotted against 2yr–100yr NOAA Atlas 14 benchmarks with alert-level coloring.

D1.5 — Technical Memo: Written description of Task 1 methodology, data sources, threshold calibration rationale, and validation results. Suitable for inclusion in a peer-reviewed methods section.

Status: D1.1 through D1.4 deliverables have been produced and validated for the Upper Sonoita Creek Watershed in early May 2026.

UPDATE — 2026-08-10 **Task 1, live.** Now produced on a 6-hour automated schedule and consumed directly by the operational dashboard, not just a one-time validation run. Real numbers, not synthetic. Gauge identity corrected — see Table 5's update box.

Task 2: LSTM training & transfer learning (Base Model → Fine-tune)

Objective: To teach the AI how desert floods work. We start with a world-class foundational model (trained on Walnut Gulch Watershed) and "fine-tune" it to our specific pilot location at Sonoita Creek. This two-stage approach allows us to launch new watersheds in a fraction of the time and with far less local data.

Step-by-Step:

- **Step 2.1 — Data Setup:** Collect historical rain and streamflow records from Walnut Gulch (our "Anchor" site) and Sonoita Creek (our "Pilot" site).
- **Step 2.2 — Base Training:** Train a "Foundation" AI model on Walnut Gulch's 70-year record. This teaches the AI the general "physics" of Arizona flash floods.
- **Step 2.3 — Benchmarking:** Compare our AI's performance against historical physics models to ensure it is accurate and balanced.
- **Step 2.4 — Local Fine-Tuning:** Take the "Anchor" model and adapt it to Sonoita Creek. Instead of retraining the whole brain, we only update the parts that need to learn local soil and terrain.
- **Step 2.5 — Event Testing:** Test the final model against major historical floods (like the 2014 Hurricane Odile remnants) to verify the timing and size of alerts.
- **Step 2.6 — Return Period Discharge Estimation:** For each watershed, derive the 2yr, 10yr, 25yr, 50yr, and 100yr return period discharge values using USGS regional regression equations for Arizona. These become the hydraulic benchmarks for Task 3.

Deliverables: D2.1 — Model Library: The Walnut Gulch "Anchor" model and the Sonoita "Pilot" model files

D2.2 — Performance Report: A simple scorecard showing how well the AI predicted flows (using metrics like NSE and Peak Timing).


D2.3 — Training Configs: The digital "recipe" files so any engineer can reproduce the exact same results later.

D2.4 — Validation Plots: Visual charts showing the AI's prediction versus actual historical flood events

D2.5 — Deployment Checklist: A short manual for moving this AI to a new watershed

UPDATE — 2026-08-10 Anchor site and architecture, both corrected (see Section 4.4's full update box above for the

diagram): Babocomari River (USGS 09471000), not Walnut Gulch, changed 2026-05-27. The "base model" is a hurdle model (LSTM event-gate + XGBoost magnitude regressor), not a single LSTM. Real held-out metrics: AUC-ROC 0.959, AUC-PR 0.643, F1 0.611, NSE 0.348, PBIAS -2.9%.



D2.5 (Deployment Checklist): still not produced as a standalone document — flagged open in the August completeness review, unresolved as of this update.

Task 3: U-Net Hydraulic Surrogate Model

Objective: Train and validate a U-Net convolutional neural network that converts predicted discharge into spatial flood inundation maps, using HEC-RAS simulations as training data.

Step-by-Step:

- **Step 3.1 — Prepare Terrain Data:** Download USGS 3DEP 1-meter lidar DEM for the Upper Sonoita Creek Watershed. The 25-watershed expansion is a later activity (Task 7). The U-Net is trained exclusively on the pilot watershed using HEC-RAS simulations specific to Sonoita Creek terrain.
- **Step 3.2 — Build HEC-RAS Models:** Construct a 2D HEC-RAS hydraulic model for the Upper Sonoita Creek Watershed using the conditioned 1-meter lidar DEM and Manning's roughness values from NLCD land cover. Calibrate against available historical high-water marks from the 2006 Santa Cruz County monsoon flood and the 2014 Odile remnant event.
- **Step 3.3 — Run HEC-RAS Simulations:** Execute 50 HEC-RAS simulations for the Upper Sonoita Creek Watershed (10 discharge levels: Q2, Q5, Q10, Q25, Q50, Q100, and 4 intermediate levels; 5 replicates with varied Manning's n to capture roughness uncertainty). Expanding to 25 watersheds × 250 simulations. Starting with 50 high-quality simulations on one watershed is sufficient to train the U-Net surrogate for operational deployment.
- **Step 3.4 — Design U-Net Architecture:** Implement a ResUNet (U-Net with residual connections) with: input channels = DEM + slope + channel distance + discharge value; output channel = flood depth (feet); 4 encoder/decoder levels; skip connections at each level.
- **Step 3.5 — Train U-Net:** Train on 50 simulations (1 watershed × 10 discharge levels + replicates), validate on 25 simulations (5 watersheds × 5 discharge levels), test on 25 simulations (5 held-out watersheds × 5 discharge levels). Use Adam optimizer, combined MSE + structural similarity loss.
- **Step 3.6 — Enforce Physics Constraints:** Post-process U-Net output to enforce: (1) gravity constraint — zero depth for pixels above predicted water surface elevation; (2) monotonicity — higher discharge = equal or greater extent; (3) mass conservation — total volume check.
- **Step 3.7 — Validate Against HEC-RAS:** Compare U-Net flood maps against HEC-RAS maps on held-out watersheds using Intersection over Union (IoU), Critical Success Index (CSI), and mean absolute depth error (MAE). Target: IoU > 0.75.
- **Step 3.8 — Generate Return Period Flood Map Library:** For each watershed, generate pre-computed flood maps for Q2, Q10, Q25, Q50, and Q100 using both HEC-RAS (reference) and U-Net (operational). These become the benchmark flood maps for comparison in the alert system.

Deliverables D3.1 — Hydraulically conditioned DEMs for all 25 watersheds, documented processing workflow.

D3.2 — 250 HEC-RAS flood simulation outputs (depth rasters), organized by watershed and discharge level. This is the primary training dataset for the U-Net.

D3.3 — Trained U-Net model: `unet_hydraulic_model.pt` — PyTorch model file with documented architecture and training history.

D3.4 — Validation report: IoU, CSI, MAE for all 25 watersheds. IoU > 0.75 on held-out watersheds.

D3.5 — Return period flood map library: Pre-computed flood maps (GeoTIFF) for Q2, Q10, Q25, Q50, Q100 for all 25 watersheds — both HEC-RAS reference and U-Net operational versions.

UPDATE — 2026-08-10 **Not what's operational (repeats Section 4.5's update):** the U-Net/HEC-RAS bank described above is implemented in code but not what drives today's maps. The operational library is the analytical FEMA/USGS depth grid described in Section 4.5's update box — 8 return periods (2/5/10/25/50/100/200/500-yr) for the single Upper Sonoita Creek pilot, not 25 watersheds. D3.4's IoU > 0.75 validation has not been run — no post-event survey polygon exists yet to validate against; this is the single largest open validation gap as of this update (Section 7's update box has the detail).

Task 4: Probabilistic Risk Products

Objective: Execute the full pipeline for all 31 GFS ensemble members and generate probabilistic flood risk products suitable for emergency management decision-making.

Step-by-Step:

- **Step 4.1 — Ensemble Pipeline Execution:** Run the complete Task 1 → Task 2 → Task 3 pipeline independently for each of the 31 GFS ensemble members. This produces 31 flood depth maps for each forecast event.
- **Step 4.2 — Compute Ensemble Statistics:** For each pixel in the flood map domain, compute: median depth (50th percentile), 90th percentile depth (worst-case), probability of exceeding 0.5m depth, standard deviation (uncertainty), and ensemble agreement index.
- **Step 4.3 — Generate Probability Map:** Create a raster where each pixel value represents the probability (0–100%) that flood depth will exceed 0.5 meters. This is the primary decision-support product for emergency managers.
- **Step 4.4 — Generate Uncertainty Map:** Create a raster showing the standard deviation of predicted depth across 31 ensemble members. High uncertainty areas require more conservative decision-making.
- **Step 4.5 — Compute Time-to-Peak:** For each ensemble member, identify the time step at which peak discharge occurs at the watershed outlet. Report the ensemble mean and spread of time-to-peak as a lead time estimate for emergency managers.
- **Step 4.6 — Generate EOC Dashboard:** Compile all probabilistic products into a multi-panel dashboard suitable for display in a county Emergency Operations Center: median map, probability map, uncertainty map, ensemble hydrograph, alert timeline, and return period comparison.

Deliverables D4.1 — Probabilistic flood map set: Median, 90th percentile, probability, and uncertainty rasters (GeoTIFF) for each forecast event.

D4.2 — EOC Dashboard: task4_eoc_dashboard.png — 6-panel visualization suitable for emergency operations center display.

D4.3 — Ensemble hydrograph: Time series of predicted discharge for all 31 ensemble members, with median and uncertainty bounds highlighted.

D4.4 — Time-to-peak report: Ensemble mean and spread of time-to-peak for each forecast event, formatted for emergency manager briefings.

Task 4, live, on the percentile-propagation method (Section 4.6's update box has the full explanation). All Table-3/D4.1–D4.4 output products are real and delivered on the live dashboard: median/likely map, worst-case map, life-safety probability map ($P > 0.5m$), uncertainty map, ensemble hydrograph, time-to-peak (Kirpich + SCS lag). Refreshed on the 6-hour schedule.

Task 5: 10-Year Storm Benchmarking and Return Period Comparison

Objective: Establish a rigorous comparison framework that allows emergency managers to understand any predicted flood event in terms of its rarity and severity relative to standard engineering design storms (2-year through 100-year return periods).

Step-by-Step:

- **Step 5.1 — Define Return Period Benchmarks:** For the Upper Sonoita Creek Watershed, compile the following benchmarks from authoritative federal sources: (a) NOAA Atlas 14 precipitation frequency estimates for 2yr, 5yr, 10yr, 25yr, 50yr, 100yr at 1hr, 3hr, 6hr, 24hr durations; (b) USGS regional regression equations for peak discharge at the same return periods; (c) pre-computed HEC-RAS flood maps for each return period (from Task 3, D3.5).
- **Step 5.2 — Build Comparison Engine:** Implement a comparison module that takes any predicted rainfall, discharge, or flood map and automatically determines: which return period bracket it falls into; how it compares to the 10-year design storm (the standard engineering benchmark); what percentage of the 100-year floodplain is inundated.
- **Step 5.3 — Historical Event Validation:** Apply the full pipeline to 5 major historical Arizona flood events (1983, 1993, 2006, 2014, 2017). Compare predicted flood maps against post-event surveys, FEMA damage assessments, and USGS peak flow records. Report IoU, CSI, and return period classification accuracy.
- **Step 5.4 — Generate Benchmark Comparison Report:** For each historical event, produce a side-by-side comparison: AFFI predicted map vs. post-event survey vs. FEMA FIRM 100-year zone. Quantify the improvement in spatial accuracy over static FIRM maps.
- **Step 5.5 — Sensitivity Analysis:** Test the pipeline's sensitivity to input uncertainty by varying: GFS forecast error ($\pm 20\%$), watershed attribute uncertainty ($\pm 10\%$), DEM resolution (1m vs. 10m). Report how these uncertainties propagate to flood map uncertainty.

Deliverables D5.1 — Return period benchmark database: NOAA Atlas 14 precipitation + USGS discharge + HEC-RAS flood maps for 2yr, 5yr, 10yr, 25yr, 50yr, 100yr for all 25 watersheds.

D5.2 — Historical event validation report: AFFI performance on 5 major AZ flood events. IoU > 0.75 vs. post-event surveys.

D5.3 — Benchmark comparison figures: Side-by-side maps (AFFI vs. post-event survey vs. FEMA FIRM) for all 5 historical events.

D5.4 — Sensitivity analysis report: Quantified uncertainty propagation from input errors to flood map errors.

D5.5 — Return period classification accuracy table: For all 5 historical events, report whether AFFI correctly classified the event's return period within one bracket (e.g., predicted "25–50 year" for an actual 35-year event).

UPDATE — 2026-08-10 — includes a real, unflattering result, reported honestly D5.4 and D5.5 were actually run this

update. Results:

- **D5.4 Sensitivity:** at the 10-year design storm, $\pm 20\%$ rainfall $\rightarrow +43.1\% / -38.4\%$ discharge (nonlinear — the SCS-CN runoff response is itself nonlinear). $\pm 10\%$ basin area $\rightarrow$ exactly $\pm 10.0\%$ discharge (linear by construction). DEM-resolution sensitivity not tested — would require regenerating the flood library at a different resolution.
- **D5.5 Classification accuracy — 0 of 4 historical events** (the catalog has 4 documented events, not 5) classified within a factor of 2 of their labeled return period against the flood library's own discharge table. Root cause: all 4 catalog events have peak discharge below the library's statistically-derived 2-year threshold — a genuine, disclosed tension between the LP-III Q2 estimate and this small event catalog, worth a domain expert's review.
- The same pass also newly measured a real historical-replay depth error: **MAE 2.301 m, bias +2.301 m** (the library over-predicts depth at every one of the 4 events). Not previously quantified — disclosed here for the first time.
- **D5.2/D5.3 IoU spatial validation:** still not done — no post-event survey polygon exists in this project's data to validate against.

Full numbers: `outputs/task5/benchmark_report.json`.

Task 6: Alert System, Government API, and Decision Dashboard

Objective: Integrate all pipeline components into an operational alert system with a secure government API and an EOC-ready decision dashboard, with alert logic defined and controlled by county emergency management agencies.

Step-by-Step:

- **Step 6.1 — Define Alert Logic:** In collaboration with Santa Cruz County OEM, define the alert threshold matrix: which combination of return period classification, probability map values, and time-to-peak triggers each alert level (Advisory / Watch / Warning). Document all decision rules in a formal Alert Logic Specification document.
- **Step 6.2 — Build Alert Engine:** Implement the alert engine as a Python module that: reads the Task 4 probabilistic outputs; applies the Task 6.1 alert logic; generates a structured JSON alert packet; logs all decisions with timestamps and input values for audit purposes.
- **Step 6.3 — Build Government API:** Develop a RESTful API (FastAPI framework) that: accepts watershed ID and forecast timestamp as inputs; executes the full AFFI pipeline; returns the alert packet and flood map URLs; requires API key authentication for all government users.
- **Step 6.4 — Build EOC Dashboard:** Develop a web-based dashboard (Streamlit or Dash framework) that: displays the current alert level with color coding; shows the median and 90th percentile flood maps on an interactive map; displays the ensemble hydrograph with uncertainty bounds; shows the return period comparison table; updates automatically every 6 hours with new GFS data.
- **Step 6.5 — Operational Testing:** Deploy the system for live testing during the Arizona monsoon season. Log all forecast events, alert levels, and (where available) observed outcomes. Compute false alarm rate, missed event rate, and lead time accuracy.
- **Step 6.6 — Documentation and Training:** Produce complete system documentation: API reference guide, dashboard user manual, alert logic specification, data flow diagram, and maintenance procedures. Conduct training sessions with Santa Cruz County OEM staff.

Deliverables D6.1 — Alert Logic Specification document: Formal definition of all alert thresholds, decision rules, and override procedures. Signed off by local OEM.

D6.2 — Alert engine module: `task6_alert_engine.py` — validated Python module with full audit logging.

D6.3 — Government API: Deployed FastAPI application with authentication, documentation, and test suite. API reference guide included.

D6.4 — EOC Dashboard: Deployed Streamlit/Dash application with user manual. Demonstrated to local OEM.

D6.5 — Operational test report: Performance metrics from live monsoon season testing: false alarm rate, missed event rate, lead time accuracy.

D6.6 — Complete system documentation package: API guide, user manual, alert logic spec, data flow diagram, maintenance procedures, training materials.

UPDATE — 2026-08-10 — the deliverable that most exceeded the original plan Delivered as a full React + TypeScript /

FastAPI application — not the Streamlit/Dash dashboard described above. Live/Simulation mode toggle, MapLibre GL interactive map with per-building and per-facility popups, Google-Flood-Hub-style gauge status badge and inundation-probability heatmap, graduated per-feature severity tiers, Decision Cockpit, NWS-style bulletin generator (cites ARS 28-910), Developer-View-gated model diagnostics, a data-staleness watchdog, and a 6-hourly APScheduler cron. Docker one-command startup for API + frontend + scheduler.

D6.1 Alert Logic Specification now exists as a standalone document: docs/

D6.1_alert_logic_specification.md (written 2026-08-10 — previously the logic existed only as code and scattered prose).

Two severe bugs were found and fixed 2026-08-10 during an internal audit, both now covered by regression tests: (1) a rainfall-unit mismatch (millimeters compared against inch thresholds with no conversion) caused a **permanent false**

WARNING / "Severe flooding" alert for about a week regardless of actual weather; (2) an out-of-memory crash in the population-exposure calculation (loading an entire national-scale raster to serve one small watershed, 6.8GB peak memory) silently froze "live" forecast data on a 6-day-old snapshot. Both fixed, verified, committed — full detail in docs/

AFFI_whitepaper_2026-08-10.md's Changelog.

D6.6 documentation package: still partial — the alert-logic spec now exists, but a standalone API reference and maintenance-procedures document do not.

6. Data Sources and Technical Feasibility

All data required for AFFI development and operation, shown in Table 5, are publicly available from federal agencies at no licensing cost. Public data eliminates a major category of project risk and ensures sustainability.

6.1 Existing Progress — Technical Feasibility

Significant foundational work has already been completed as of early May 2026, demonstrating technical feasibility, capability and commitment:

1. A prototype meteorological forecast interface has been developed and tested for the Sonoita Creek watershed (HUC 150503010204), successfully ingesting GFS ensemble data and computing mean areal precipitation in real time.
2. Alert threshold logic has been calibrated against NOAA Atlas 14 precipitation frequency data for Santa Cruz County, with Advisory / Watch / Warning levels benchmarked against 10-year, 25-year, and 50-year return periods.
3. The complete data pipeline architecture has been designed, documented, and partially implemented.
4. A dummy rainfall stress-test framework has been developed and validated against NOAA Atlas 14 IDF benchmarks for all return periods from 2-year through 100-year.
5. Task 1 deliverables D1.1 through D1.4 are complete and validated.

Table 5: Data Sources for AFFI

Data Type	Source	Resolution	Coverage	Cost
Streamflow records (anchor)	USDA ARS — Walnut Gulch Experimental Watershed	Sub-hourly	0+ years, 89 rain gauges, 12 flumes	Free
Historical precipitation	ERA5 via Open-Meteo Archive API	0.25°, hourly	1940–present, global	Free
Streamflow (pilot watershed)	USGS NWIS — Gauge 09481740 (Santa Cruz at Tubac)	15-min interval	1995–present, 30yr record	Free
Terrain (DEM)	USGS 3DEP	1-meter lidar	Statewide AZ	Free
Real-time forecast	NOAA GFS	~13km, 6-hourly	31 ensemble members	Free
High-res forecast	NOAA HRRR	3km, hourly	18 ensemble members	Free
Radar rainfall	NOAA MRMS	1km, 2-min	Real-time, CONUS	Free
Soil data	USDA SSURGO	1:24,000	Statewide AZ	Free
Land cover	USGS NLCD	30m	Statewide AZ	Free
Watershed attributes	USGS GAGES-II	Per watershed	9,322 US watersheds	Free
Flood benchmarks	NOAA Atlas 14	Per county	All AZ counties	Free

UPDATE — 2026-08-10

- **Anchor streamflow row:** Babocomari River (USGS 09471000), real gauge, not Walnut Gulch ("0+ years" was also a likely typo in the original for "70+ years" — moot now that Walnut Gulch isn't the anchor).
- **Pilot watershed gauge — corrected:** USGS **09481500** (Sonoita Creek near Patagonia), not 09481740 (Santa Cruz at Tubac). 09481740 does not directly instrument this watershed's own outflow; 09481500 does, and is what the code, config, and live API all actually use.
- **DEM resolution:** operational library runs on the 10-m USGS 3DEP product, not 1-m lidar (1-m is available and is the target for full HEC-RAS work).
- **Two rows now missing that should be added:** FEMA NFHL + FEMA FIS Base Flood Elevations (regulatory floodplain + BFE, Free, nationwide where mapped) and USGS Log-Pearson III discharge estimates (statewide AZ, Free) — both now drive the operational Stage 3 library and aren't listed in the original table above.

7. Validation Plan and Performance Goals

AFFI will be validated against four independent benchmarks before full operational certification. Specific, measurable performance goals (targets) are stated below in Table 6. These targets are based on published standards in the peer-reviewed hydrology literature.

Table 6: Performance Goals

Validation Level	Method	Target	Reference Standard
In-sample accuracy	LSTM tested on held-out years (2001–2006) from 20 training watersheds	NSE > 0.75, KGE > 0.70	Moriasi et al. 2007
Out-of-sample generalization	LSTM tested on 5 completely held-out watersheds never seen in training	NSE > 0.65	Kratzert et al. 2019
Historical event validation	Full pipeline tested on 5 major AZ flood events (1983, 1993, 2006, 2014, 2017)	IoU > 0.75 vs. post-event surveys	FEMA post-disaster assessments
Speed benchmark	Wall-clock time from GFS input to flood map output	< 60 seconds	Flash flood warning lead time requirement

UPDATE — 2026-08-10 **Targets retained as targets; here is what was actually achieved, honestly:**

- **Hydrology (Upper Sonoita Creek, held-out test):** AUC-ROC 0.959, AUC-PR 0.643, F1 0.611 (at $P \geq 0.85$), **NSE 0.348** (below the 0.75 target), PBIAS -2.9% . The NSE gap is a known property of arid, ephemeral basins — a metric normalized by observed variance is punished hardest exactly where flow is near-zero most of the time (Kratzert et al. 2019 finds arid basins the hardest rainfall-runoff regime, lowest-NSE class). Detection (AUC-ROC), not magnitude, drives the operational go/no-go decision.
- **Speed:** ~60 seconds end-to-end — target met.
- **Historical event / IoU validation:** not done — see Task 5's update box. No post-event survey polygon exists to validate against; this is the single largest remaining validation gap.
- **Operational integrity incident (disclosed):** from ~August 4–10, 2026, a unit-conversion bug caused a week of false WARNING alerts, and a separate memory bug froze live data on a 6-day-old snapshot. Both found, fixed, and disclosed — see Task 6's update box.

7.1 Baseline Comparisons

AFFI will be compared against three baselines to demonstrate its contribution:

- **Baseline 1 — Simple LSTM (no watershed attributes):** Proves that physics-guided design adds value over generic AI
- **Baseline 2 — HEC-HMS + HEC-RAS (full physics):** Proves that AFFI approaches physics-model accuracy at 360× the speed
- **Baseline 3 — FEMA static FIRM maps:** Proves that dynamic event-based maps outperform static hazard maps for real-time decisions

8. Transferability to Other Watersheds

Design Principle: Every component of AFFI was designed from the ground up to be transferable to any watershed in Arizona or the United States. The Upper Sonoita Creek Watershed is the pilot but not the only watershed that could be used. The AFFI architecture is transferable by design.

To deploy AFFI for a new watershed, the following watershed-specific inputs must be substituted. All other components (LSTM architecture, U-Net architecture, alert logic framework, API, dashboard) remain unchanged:

Component	What Changes Per Watershed	Data Source	Estimated Effort
AOI / Bounding Box	Watershed boundary polygon and bounding box coordinates	USGS National Map	1 hour
NOAA Atlas 14 Benchmarks	County-specific precipitation frequency estimates	NOAA Atlas 14	2 hours
Alert Thresholds	Advisory/Watch/Warning rainfall thresholds calibrated to local conditions	County OEM + NOAA Atlas 14	4 hours
DEM	1-meter lidar DEM for new watershed domain	USGS 3DEP	4 hours
HEC-RAS Model	New 2D hydraulic model for new watershed	DEM + NLCD + field survey	20–40 hours
LSTM Fine-Tuning	Optional: fine-tune on local gauge data if available	USGS NWIS	8 hours
U-Net Re-Training	Add new HEC-RAS simulations to training set; retrain U-Net	HEC-RAS outputs from above	10 hours

Estimated total effort per new watershed: 40–80 hours, depending on whether local gauge data is available for LSTM fine-tuning. This is the marginal cost of statewide expansion — not the full system development cost.

Watersheds where AFFI is immediately applicable without modification: Any Arizona watershed with an available USGS 3DEP lidar DEM and NOAA Atlas 14 county data. For watersheds outside Arizona, the only additional requirement is downloading the relevant NOAA Atlas 14 volume for the applicable state.

UPDATE — 2026-08-10 Still one watershed validated to date; the architecture is engineered so the next one is a data-substitution exercise, not a rebuild. A key transferability finding since May: because the operational Stage-3 library is built from nationally available FEMA and USGS open data (not bespoke HEC-RAS runs), any watershed with FEMA NFHL AE mapping, published BFEs, and a 3DEP DEM can receive an operational flood library quickly — the analytical method turned out to be a transferability asset in its own right, not just a pilot shortcut.

Statewide Scaling — Transfer to Additional Arizona Watersheds

Objective: Deploy the validated AFFI pipeline to additional Arizona watersheds using the transfer learning framework established in Tasks 1–6, demonstrating that the marginal cost of adding a new watershed is 40–80 hours of engineering effort, not a full system rebuild.

Key steps:

- **Step 8.1.1 — Watershed Selection:** Prioritize AZ watersheds using 4 criteria: (a) USGS 3DEP 1-meter lidar available, (b) NOAA Atlas 14 county data available, (c) flash flood risk to populated area, (d) no existing operational flood warning system.
- **Step 8.1.2 — Rapid Deployment Protocol:** For each new watershed, execute the 8-component checklist from Section 8 (Scalability). Target: full deployment in 40–80 hours per watershed.
- **Step 8.1.3 — Cross-Watershed Validation:** Test whether the Sonoita Creek fine-tuned LSTM (Task 2b) transfers to a new watershed without retraining. Report NSE degradation vs. full fine-tuning. This directly tests generalizability.
- **Step 8.1.4 — Statewide Dashboard:** Extend the Task 6 EOC Dashboard to display alert levels for all deployed watersheds on a single Arizona-wide map.

Deliverables: D8.1.1 — Rapid deployment checklist: Documented per-watershed configuration steps, verified on 3 pilot Arizona watersheds beyond the Sonoita Creek pilot project.

D8.1.2 — Cross-watershed generalization report: NSE, KGE on 3 new watersheds with and without fine-tuning.

D8.1.3 — Statewide alert dashboard: Extended Streamlit/Dash application showing all deployed watersheds.

D8.1.4 — Deployment cost audit: Actual engineering hours per watershed for the 3 pilot expansions.

UPDATE — 2026-08-10 Correctly still a roadmap item, not falsely claimed done — `config/watersheds/upper_sonoita.yaml` is the only watershed configuration that exists today. No change to report here.

9. References (Verified Hyperlinks)

1. Cloke, H. L., & Pappenberger, F. (2009). Ensemble flood forecasting: A review. *Journal of Hydrology*. doi.org/10.1016/j.jhydrol.2009.06.005
2. Federal Emergency Management Agency (FEMA). (2022). Flood Insurance Rate Maps (FIRMs). fema.gov/flood-maps
3. IPCC. (2021). Climate Change 2021: The Physical Science Basis. Sixth Assessment Report. ipcc.ch/report/ar6/wg1
4. Kratzert, F., et al. (2019). Towards learning universal, regional, and local hydrological behaviors via machine learning. *Nature Communications*. doi.org/10.1038/s41467-019-13507-2
5. Moriasi, D. N., et al. (2007). Model evaluation guidelines for systematic quantification of accuracy in watershed simulations. *Transactions of the ASABE*. doi.org/10.13031/2013.23153
6. NOAA Office of Water Prediction. (2024). National Water Model (NWM) Documentation. water.noaa.gov/about/nwm
7. NOAA Hydrometeorological Design Studies Center. (2024). Precipitation Frequency Data Server (Atlas 14). hdsc.nws.noaa.gov/hdsc/pfds
8. Ronneberger, O., et al. (2015). U-Net: Convolutional networks for biomedical image segmentation. *MICCAI*. doi.org/10.1007/978-3-319-24574-4_28
9. Sit, M., et al. (2020). A comprehensive review of deep learning applications in hydrology. *Water Science and Technology*. doi.org/10.2166/wst.2020.369
10. USGS. (2024). 3D Elevation Program (3DEP) High-Resolution Lidar Data. usgs.gov/3d-elevation-program
11. USGS. (2011). Geospatial Attributes of Gages for Evaluating Streamflow (GAGES-II). water.usgs.gov/GIS/metadata
12. USGS National Water Information System (NWIS). (2024). Real-time Streamflow Data for Arizona. waterdata.usgs.gov/nwis
13. Tellman, B., et al. (2021). Satellite imaging reveals increased proportion of population exposed to floods. *Nature*. doi.org/10.1038/s41586-021-03695-w
14. Google Research. (2024). Flood Hub Global Flood Forecasting. sites.research.google/floods
15. Multi-Resolution Land Characteristics (MRLC) Consortium. (2021). National Land Cover Database (NLCD). mrlc.gov/data
16. Oregon State University. (2024). PRISM Climate Group Precipitation Data. prism.oregonstate.edu
17. NOAA National Severe Storms Laboratory. (2024). Multi-Radar Multi-Sensor (MRMS) System. mrms.nssl.noaa.gov
18. USDA Natural Resources Conservation Service. (2024). Soil Survey Geographic Database (SSURGO). nrcs.usda.gov/ssurgo

UPDATE — 2026-08-10 — additional references now relevant Methods actually in the deployed code that weren't cited in

the original reference list:

1. Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *KDD*. doi.org/10.1145/2939672.2939785
2. Lin, T.-Y., et al. (2017). Focal loss for dense object detection. *ICCV*. doi.org/10.1109/ICCV.2017.324
3. USDA Natural Resources Conservation Service. (1986). Urban Hydrology for Small Watersheds (TR-55, SCS Curve Number method).
4. Leopold, L. B., & Maddock, T. (1953). The hydraulic geometry of stream channels and some physiographic implications. USGS Professional Paper 252.
5. Kirpich, Z. P. (1940). Time of concentration of small agricultural watersheds. *Civil Engineering*, 10(6), 362.
6. Federal Emergency Management Agency (FEMA). National Flood Hazard Layer (NFHL). fema.gov/flood-maps/national-flood-hazard-layer

Where to find the complete August 2026 detail

This document preserves the original May 11, 2026 whitepaper exactly, with dated update boxes added at each point where the deployed system has since diverged. For the full, continuously-maintained technical writeup — including everything in the update boxes above written out in complete prose, plus dashboard features built since May that have no May-11 counterpart to annotate — see:

- `docs/AFFI_whitepaper_2026-08-10.md / .pdf` — the current full whitepaper
- `docs/D6.1_alert_logic_specification.md` — standalone alert-logic reference
- `STATUS.md` — running project status log